

# machine\_learning\_project

*Subash Pathak*

*April 09, 2018*

## Practical Machine Learning Assignment

### Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

### Data

The training data for this project are available here: [<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>]

The test data are available here: [<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>]

The data for this project come from this source: [<http://groupware.les.inf.puc-rio.br/har>]. If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment.

### What you should submit

The goal of your project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

Your submission should consist of a link to a Github repo with your R markdown and compiled HTML file describing your analysis. Please constrain the text of the writeup to < 2000 words and the number of figures to be less than 5. It will make it easier for the graders if you submit a repo with a gh-pages branch so the HTML page can be viewed online (and you always want to make it easy on graders :-). You should also apply your machine learning algorithm to the 20 test cases available in the test data above. Please submit your predictions in appropriate format to the programming assignment for automated grading. See the programming assignment for additional details.

## Approach used in the analysis:

Our outcome variable is `classe`, a factor variable. For this data set, “participants were asked to perform one set of 10 repetitions of the Unilateral Dumbbell Biceps Curl in 5 different fashions: - exactly according to the specification (Class A) - throwing the elbows to the front (Class B) - lifting the dumbbell only halfway (Class C) - lowering the dumbbell only halfway (Class D) - throwing the hips to the front (Class E)

Two models will be tested using decision tree and random forest. The model with the highest accuracy will be chosen as our final model.

## Cross-validation

Cross-validation will be performed by subsampling our training data set randomly without replacement into 2 subsamples: `TrainTrainingSet` data (70% of the original Training data set) and `TestTrainingSet` data (30%). Our models will be fitted on the `TrainTrainingSet` data set, and tested on the `TestTrainingSet` data. Once the most accurate model is chosen, it will be tested on the original Testing data set.

## Expected out-of-sample error

The expected out-of-sample error will correspond to the quantity: 1-accuracy in the cross-validation data. Accuracy is the proportion of correct classified observation over the total sample in the `TestTrainingSet` data set. Expected accuracy is the expected accuracy in the out-of-sample data set (i.e. original testing data set). Thus, the expected value of the out-of-sample error will correspond to the expected number of misclassified observations/total observations in the Test data set, which is the quantity: 1-accuracy found from the cross-validation data set.

Our outcome variable “`classe`” is a factor variable. We split the Training dataset into `TrainTrainingSet` and `TestTrainingSet` datasets.

## Packages, Libraries and Seed

Installing packages, loading libraries, and setting the seed for reproducibility:

```
rm(list=ls())  
library(knitr)
```

```
## Warning: package 'knitr' was built under R version 3.3.3
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.3.3
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 3.3.2
```

```
library(rpart)
```

```
## Warning: package 'rpart' was built under R version 3.3.3
```

```
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 3.3.3
```

```
library(rattle)
```

```
## Warning: package 'rattle' was built under R version 3.3.3
```

```
## Rattle: A free graphical interface for data science with R.
```

```
## Version 5.1.0 Copyright (c) 2006-2017 Togaware Pty Ltd.
```

```
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.3.3
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:rattle':
```

```
##
```

```
##     importance
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##     margin
```

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 3.3.3
```

```
## corrplot 0.84 loaded
```

## cleaning the data

```
set.seed(2468)
```

```
trainUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
```

```
testUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"
```

```
training <- read.csv(url(trainUrl), na.strings=c("NA", "#DIV/0!", ""))
```

```
testing <- read.csv(url(testUrl), na.strings=c("NA", "#DIV/0!", ""))
```

```
#delete columns with all missing values
```

```
trainingset<-training[,colSums(is.na(training)) == 0]
```

```
testingset <-testing[,colSums(is.na(testing)) == 0]
```

```
# Delete variables are not applicable to our current project: user_name, raw_timestamp_part_1, raw_time
```

```
trainingset <-trainingset[,-c(1:7)]
```

```
testingset <-testingset[,-c(1:7)]
```

```
# partition the data so that 70% of the training dataset into training and the remaining 30% to testing
```

```
traintrainset <- createDataPartition(y=trainingset$classe, p=0.70, list=FALSE)
```

```
TrainTrainingSet <- trainingset[traintrainset, ]
```

```
TestTrainingSet <- trainingset[-traintrainset, ]
```

```
# The variable "classe" contains 5 levels: A, B, C, D and E. A plot of the outcome variable will allow
```

```
plot(TrainTrainingSet$classe, col="red", main="Plot of levels of variable classe within the TrainTrainingSet data set")
```

## Plot of levels of variable classe within the TrainTrainingSet data set



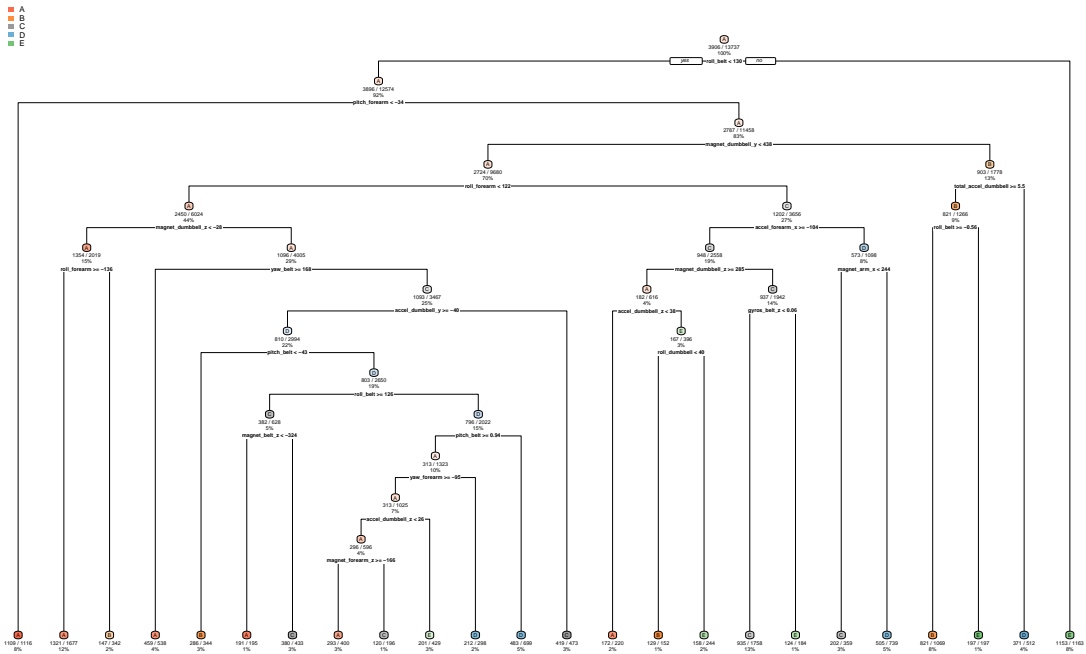
From the graph above level A is the most frequent while D is the least frequent

## Prediction Model 1 Decision tree

```
model1 <- rpart(classe ~ ., data=TrainTrainingSet, method="class")
prediction1 <- predict(model1, TestTrainingSet, type = "class")

# Plot the Decision Tree
rpart.plot(model1, main="Classification Tree", extra=102, under=TRUE, faclen=0)
```

## Classification Tree



```
# Test results on our TestTrainingSet data set:
confusionMatrix(prediction1, TestTrainingSet$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1505  176   22   53   19
##           B   54  601   56   71   79
##           C   39  214  888  102  134
##           D   59   68   58  647   75
##           E   17   80    2   91  775
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.7504
##           95% CI : (0.7391, 0.7614)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
##           Kappa : 0.6839
##           McNemar's Test P-Value : < 2.2e-16
```

```
## Statistics by Class:
```

```
##
##           Class: A Class: B Class: C Class: D Class: E
```

## Sensitivity	0.8990	0.5277	0.8655	0.6712	0.7163
## Specificity	0.9359	0.9452	0.8994	0.9472	0.9604
## Pos Pred Value	0.8479	0.6980	0.6449	0.7133	0.8031
## Neg Pred Value	0.9589	0.8929	0.9694	0.9363	0.9376
## Prevalence	0.2845	0.1935	0.1743	0.1638	0.1839
## Detection Rate	0.2557	0.1021	0.1509	0.1099	0.1317
## Detection Prevalence	0.3016	0.1463	0.2340	0.1541	0.1640
## Balanced Accuracy	0.9175	0.7364	0.8824	0.8092	0.8384

## Prediction model 2: Random Forest

```
model2 <- randomForest(classe ~. , data=TrainTrainingSet, method="class")

# Predicting:
prediction2 <- predict(model2, TestTrainingSet, type = "class")

# Test results on TestTrainingSet data set:
confusionMatrix(prediction2, TestTrainingSet$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1673    6    0    0    0
##           B    1 1132    7    0    0
##           C    0    1 1019   12    0
##           D    0    0    0  952    0
##           E    0    0    0    0 1082
##
## Overall Statistics
##
##           Accuracy : 0.9954
##           95% CI : (0.9933, 0.997)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9942
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9994  0.9939  0.9932  0.9876  1.0000
## Specificity      0.9986  0.9983  0.9973  1.0000  1.0000
## Pos Pred Value   0.9964  0.9930  0.9874  1.0000  1.0000
## Neg Pred Value   0.9998  0.9985  0.9986  0.9976  1.0000
## Prevalence       0.2845  0.1935  0.1743  0.1638  0.1839
## Detection Rate   0.2843  0.1924  0.1732  0.1618  0.1839
## Detection Prevalence 0.2853  0.1937  0.1754  0.1618  0.1839
## Balanced Accuracy 0.9990  0.9961  0.9953  0.9938  1.0000
```

## Decison on which model to use

Random forest performed better than decision tree model. We have accuracy 0.9954(95% CI: 0.9933,0.997) for random forest method while decision tree method has accuracy 0.7504(95% CI: 0.7391,0.7614). The expected out of sample error is estimated at 0.005 (5%)

## Submission

Final outcome based on the prediction model2 (random forest) applied against the original testing set:

```
final <- predict(model2, testingset, type="class")
final
```

```
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E
```